1 Neural ODE

1.1 Adjoint Method

Let M a manifold, and $x(t) \in C^1(\mathbb{R}, M)$ a trajectory, obeying

$$\frac{\mathrm{d}x}{\mathrm{d}t}(t) = f(t, x(t); \theta),$$

and

$$x(t_0) = x_0,$$

where $f \in C(\mathbb{R} \times M, T_M)$ parameterized by θ. For $\forall t_1 > t_0$, let

$$x_1 := x_0 + \int_{t_0}^{t_1} f(t, x(t); \theta) \mathrm{d}t.$$

Then

Theorem 1. *Let $\mathcal{C} \in C^1(M, \mathbb{R})$, and $\forall x(t) \in C^1(\mathbb{R}, M)$ obeying dynamics $f(t, x; \theta) \in C^1(\mathbb{R} \times M, T_M)$ with initial value $x(t_0) = x_0$. Denote*

$$L := \mathcal{C}\left(x_0 + \int_{t_0}^{t_1} f(\tau, x(\tau); \theta) \mathrm{d}\tau \right).$$

Then we have, for $\forall t \in [t_0, t_1]$ given,

$$\frac{\partial L}{\partial x^\alpha(t)} = \frac{\partial L}{\partial x_1^\alpha} - \int_t^{t_1} \frac{\partial L}{\partial x^\beta(\tau)} \frac{\partial f^\beta}{\partial x^\alpha}(\tau, x(\tau); \theta) \mathrm{d}\tau,$$

and

$$\frac{\partial L}{\partial \theta} = - \int_{t_0}^{t_1} \frac{\partial L}{\partial x^\beta(\tau)} \frac{\partial f^\beta}{\partial \theta}(\tau, x(\tau); \theta) \mathrm{d}\tau.$$

Proof. Suppose the $x(t)$ is layerized, the L depends on the variables (inputs and model parameters) on the ith layer can be regarded as the loss of a new model by truncating the original at the ith layer, which we call $L_i(z_i)$.

$$\begin{aligned} \frac{\partial L_i}{\partial x_i^\alpha}(x_i) &= \frac{\partial L_{i+1}}{\partial x_{i+1}^\beta}(x_{i+1}) \frac{\partial x_{i+1}^\beta}{\partial x_i^\alpha}(x_i) \\ &= \frac{\partial L_{i+1}}{\partial x_1^\beta}(x_{i+1}) \frac{\partial}{\partial x_i^\alpha}(x_i^\beta + f^\beta(t_i, x_i; \theta) \Delta t) \\ &= \frac{\partial L_{i+1}}{\partial x_{i+1}^\alpha}(x_{i+1}) + \frac{\partial L_{i+1}}{\partial x_{i+1}^\beta}(x_{i+1}) \partial_\alpha f^\beta(t_i, x_i; \theta) \Delta t. \end{aligned}$$

This hints that

$$\frac{\mathrm{d}}{\mathrm{d}t} \frac{\partial L}{\partial x^\alpha(t)} = - \frac{\partial L}{\partial x^\beta(t)} \frac{\partial f^\beta}{\partial x^\alpha(t)}(t, x(t); \theta).$$

The initial value is $\partial L / \partial x_1$. Thus

$$\frac{\partial L}{\partial x(t)} = \frac{\partial L}{\partial x_1} - \int_t^{t_1} \frac{\partial L}{\partial x^\beta(\tau)} \frac{\partial f^\beta}{\partial x^\alpha(\tau)}(\tau, x(\tau); \theta) \mathrm{d}\tau.$$

Varing θ will vary the $L_i(x_i)$ from two aspects, the effect from $\partial L_{i+1}/\partial\theta$ and the Δx_{i+1} caused by $\Delta\theta$.

$$\begin{aligned}
\frac{\partial L_i}{\partial\theta}(x_i) &= \frac{\partial L_{i+1}}{\partial\theta}(x_{i+1}) + \frac{\partial L_{i+1}}{\partial x_{i+1}}\frac{\partial x_{i+1}}{\partial\theta}\\
&= \frac{\partial L_{i+1}}{\partial\theta}(x_{i+1}) + \frac{\partial L_{i+1}}{\partial x_{i+1}}\frac{\partial}{\partial\theta}(x_i^\beta + f^\beta(t_i,x_i;\theta)\Delta t)\\
&= \frac{\partial L_{i+1}}{\partial\theta}(x_{i+1}) + \frac{\partial L_{i+1}}{\partial x_{i+1}}\frac{\partial f^\beta}{\partial\theta}(t_i,x_i;\theta)\Delta t.
\end{aligned}$$

This hints that

$$\frac{\mathrm{d}}{\mathrm{d}t}\frac{\partial L}{\partial\theta} = -\frac{\partial L}{\partial x^\alpha(t)}\frac{\partial f^\beta}{\partial\theta}(t,x(t),\theta).$$

The initial value is 0 since $\mathcal{C}(.)$ is explicitly independent on θ. Thus

$$\frac{\partial L}{\partial\theta} = -\int_{t_0}^t \frac{\partial L}{\partial x^\beta(\tau)}\frac{\partial f^\beta}{\partial\theta}(\tau,x(\tau);\theta)\mathrm{d}\tau. \qquad\qquad \square$$

2 Hopfield Network

2.1 Discrete-time Hopfield Network

Definition 2. *[Discrete-time Hopfield Network]*

Let $t \in \mathbb{N}$ and $x \in \{-1,+1\}^d$, $W \in \mathbb{R}^d \times \mathbb{R}^d$ with $W_{\alpha\beta} = W_{\beta\alpha}$ and $W_{\alpha\alpha} = 0$, and $b \in \mathbb{R}^d$. Define discrete-time dynamics

$$x^\alpha(t+1) = \mathrm{sign}(W^\alpha{}_\beta x^\beta(t) + b^\alpha).$$

The (x,W,b) is called a discrete-time Hopfield network.

Lemma 3. Let (x,W,b) a discrete-time Hopfield network. Define $\mathcal{E}(x) := -(1/2)W_{\alpha\beta}x^\alpha x^\beta - b_\alpha x^\alpha$. Then $\mathcal{E}(x(t+1)) - \mathcal{E}(x(t)) \leqslant 0$.

Proof. Consider async-updation of Hopfield network, that is, change the component at dimension $\hat{\alpha}$, i.e. $x'_{\hat{\alpha}} = \mathrm{sign}[W_{\hat{\alpha}\beta}x^\beta + b_{\hat{\alpha}}]$, then

$$\begin{aligned}
\mathcal{E}(x') - \mathcal{E}(x) &= -\frac{1}{2}W_{\alpha\beta}x'^\alpha x'^\beta - b_\alpha x'^\alpha + \frac{1}{2}W_{\alpha\beta}x^\alpha x^\beta + b_\alpha x^\alpha\\
&= -2\left(x'^{\hat{\alpha}} - x^{\hat{\alpha}}\right)\left(W_{\hat{\alpha}\beta}x^\beta + b_{\hat{\alpha}}\right),
\end{aligned}$$

which employs conditions $W_{\alpha\beta} = W_{\beta\alpha}$ and $W_{\alpha\alpha} = 0$. Next, we prove that, combining with $x'_{\hat{\alpha}} = \mathrm{sign}[W_{\hat{\alpha}\beta}x^\beta + b_{\hat{\alpha}}]$, this implies $\mathcal{E}(x') - \mathcal{E}(x) \leqslant 0$.

If $(x'^{\hat{\alpha}} - x^{\hat{\alpha}}) > 0$, then $x'^{\hat{\alpha}} = 1$ and $x^{\hat{\alpha}} = -1$. Since $x'_{\hat{\alpha}} = \mathrm{sign}[W_{\hat{\alpha}\beta}x^\beta + b_{\hat{\alpha}}]$, $W_{\hat{\alpha}\beta}x^\beta + b_{\hat{\alpha}} > 0$. Then $\mathcal{E}(x') - \mathcal{E}(x) < 0$. Contrarily, if $(x'^{\hat{\alpha}} - x^{\hat{\alpha}}) < 0$, then $x'^{\hat{\alpha}} = -1$ and $x^{\hat{\alpha}} = 1$, implying $W_{\hat{\alpha}\beta}x^\beta + b_{\hat{\alpha}} < 0$. Also $\mathcal{E}(x') - \mathcal{E}(x) < 0$. Otherwise, $\mathcal{E}(x') - \mathcal{E}(x) = 0$. So, we conclude $\mathcal{E}(x') - \mathcal{E}(x) \leqslant 0$. $\qquad\square$

Theorem 4. *[Convergene of Discrete-time Hopfield Network] Let (x, W, b) a discrete-time Hopfield network. Then any trajectory obeying the update rule will converge either to a fixed point or a limit circle.*

Proof. Since the states of the network are finite, the $\mathcal{E}$ is lower bounded. $\square$

2.2 Continuous-time Hopfield Network

Definition 5. *[Continuous-time Hopfield Network]*

Let $t \in [0, +\infty)$ and $x \in [-1, +1]^d$, $W \in \mathbb{R}^d \times \mathbb{R}^d$ with $W_{\alpha\beta} = W_{\beta\alpha}$, and $b \in \mathbb{R}^d$. Define dynamics

$$\tau \frac{\mathrm{d}x^\alpha}{\mathrm{d}t}(t) = -x^\alpha(t) + f(W^\alpha{}_\beta x^\beta(t) + b^\alpha),$$

where $\tau \in (0, +\infty)$ a constant and $f \colon \mathbb{R} \to [-1, 1]$ being increasing. The $(x, W, b; \tau, f)$ is called a continuous-time Hopfield network.

Remark 6. With

$$\tau \frac{x^\alpha(t + \Delta t) - x^\alpha(t)}{\Delta t} = -x^\alpha(t) + f(W^\alpha{}_\beta x^\beta(t) + b^\alpha)).$$

Setting $\Delta t = \tau$ gives and $f(.) = \mathrm{sign}(.)$ gives

$$x^\alpha(t + \tau) = \mathrm{sign}(W^\alpha{}_\beta x^\beta(t) + b^\alpha),$$

which is the same as the discrete-time Hopfield network.

Lemma 7. *Let $(x, W, b; \tau, f)$ a continous-time Hopfield network. Define $a^\alpha := W^\alpha{}_\beta x^\beta + b^\alpha$ and $y^\alpha := f(a^\alpha)$, then*

$$\mathcal{E}(y) := -\frac{1}{2} W_{\alpha\beta} y^\alpha y^\beta - b_\alpha y^\alpha + \sum_\alpha \int^{y^\alpha} f^{-1}(y^\alpha) \mathrm{d}y^\alpha.$$

Then $\mathcal{E}(y(x(t + \mathrm{d}t))) - \mathcal{E}(y(x(t))) \leqslant 0$.

Proof. The dynamics of a^α is

$$\begin{aligned}
\tau \frac{\mathrm{d}a^\alpha}{\mathrm{d}t} &= \tau W^\alpha{}_\beta \frac{\mathrm{d}x^\beta}{\mathrm{d}t} \\
&= W^\alpha{}_\beta [-x^\beta(t) + f(a^\beta)] \\
&= -(W^\alpha{}_\beta x^\beta(t) + b^\alpha) + b^\alpha + W^\alpha{}_\beta y^\beta \\
&= W^\alpha{}_\beta y^\beta + b^\alpha - a^\alpha.
\end{aligned}$$

Since W is symmetric, we have $\partial \mathcal{E} / \partial y^\alpha = -W_{\alpha\beta} y^\beta - b_\alpha + f^{-1}(y_\alpha)$. Then

$$\begin{aligned}
\frac{\mathrm{d}\mathcal{E}}{\mathrm{d}t} &= \frac{\mathrm{d}y^\alpha}{\mathrm{d}t}(-W_{\alpha\beta} y^\beta - b_\alpha + f^{-1}(y_\alpha)) \\
&= \frac{\mathrm{d}y^\alpha}{\mathrm{d}t}(-W_{\alpha\beta} y^\beta - b_\alpha + a_\alpha) \\
&= -\frac{\mathrm{d}y^\alpha}{\mathrm{d}t}(W_{\alpha\beta} y^\beta + b_\alpha - a_\alpha)
\end{aligned}$$

Notice that, the second term of rhs is exactly the dynamics of a_α, then

$$
\begin{aligned}
\frac{\mathrm{d}\mathcal{E}}{\mathrm{d}t} &= -\tau \frac{\mathrm{d}y^\alpha}{\mathrm{d}t} \frac{\mathrm{d}a_\alpha}{\mathrm{d}t} \\
&= -\tau \frac{\mathrm{d}y^\alpha}{\mathrm{d}a^\alpha} \left(\frac{\mathrm{d}a^\alpha}{\mathrm{d}t} \frac{\mathrm{d}a_\alpha}{\mathrm{d}t} \right) \\
&= -\tau f'(a^\alpha) \left(\frac{\mathrm{d}a^\alpha}{\mathrm{d}t} \frac{\mathrm{d}a_\alpha}{\mathrm{d}t} \right).
\end{aligned}
$$

Since f is increasing and $\tau > 0$, $\mathrm{d}\mathcal{E}/\mathrm{d}t \leqslant 0$. $\qquad\square$

Remark 8. The condition $W_{\alpha\alpha} = 0$ for $\forall \alpha$ is not essential for this lemma. Indeed, this condition is absent in the proof. This differs from the case of discrete-time.

Theorem 9. *[Convergene of Continuous-time Hopfield Network] Let $(x, W, b; \tau, f)$ a continous-time Hopfield network. Then any trajectory along the dynamics will converge either to a fixed point or a limit circle.*

Proof. The function $E := \mathcal{E} \circ y$ is lower bounded since y, i.e. function $f: \mathbb{R} \to [-1, 1]$, is bounded. This E is a Lyapunov function for the continous-time Hopfield network. $\qquad\square$

Corollary 10. *Let $(x, W, b; \tau, f)$ a continous-time Hopfield network. And $D := \{x_n | x_n \in \mathbb{R}^d, n = 1, ..., N\}$ a dataset[1]. If add constraint $W_{\alpha\alpha} = 0$ for $\forall \alpha$, then we can train the Hopfield nework by seeking a proper parameters (W, b), s.t. its stable points cover the dataset as much as possible, by[2]*

Algorithm **1**

```
W, b = init_W, init_b  # e.g. by Glorot initializer
for step in range(max_step):
    for x in dataset:
        y = f(W @ x + b)
        loss = norm(x - y)
        optimizer.minimize(objective=loss, variables=(W, b))
        W = set_zero_diag(symmetrize(W))
```

Proof. For $\forall x_n \in D$, we try to find (W, b), s.t. $\mathrm{d}x/\mathrm{d}t = 0$ at x_n, i.e.

$$
x_n^\alpha = f(W^\alpha{}_\beta x_n^\beta + b^\alpha).
$$

When $W_{\alpha\alpha} = 0$ for $\forall \alpha$, $f(W^\alpha{}_\beta x^\beta + b^\alpha)$ thus has no information of x^α, it has to predict the x^α by the interaction between x^α and the other x's components. $\qquad\square$

Remark 11. This algorithm is equivalent to

Algorithm 2

```
dt = ...  # e.g. 0.1
W, b = init_W, init_b
for step in range(max_step):
    for x in dataset:
        # that is, compute x(dt), with x(0) = x
```

---

1. We use Greek alphabet for component in $\mathbb{R}^d$ and Lattin alphabet for element in dataset.
2. This algorithm generalizes the algorithm 42.9 of Mackay.

```
        y = ode_solve(f=lambda t, x: -x + f(W @ x + b), t0=0, t1=dt, x0=x)
        loss = norm(x - y)
        optimizer.minimize(objective=loss, variables=(W, b))
        W = set_zero_diag(symmetrize(W))
```

Indeed, trying to reach $y = x$ within a small interval will force x to be a fixed point.

3 Variations

3.1 Variation 1

Theorem 12. *Let $v \in \mathbb{R}^d$, $F \in C^1(\mathbb{R}^n, \mathbb{R})$, $W \in \mathbb{R}^n \times \mathbb{R}^d$, $b \in \mathbb{R}^n$, and $\tau > 0$. Define the dynamics*

$$\tau \frac{\mathrm{d}x}{\mathrm{d}t} = -\nabla E(x) = -x + W^T \cdot \nabla F(W \cdot x + b) + v.$$

If $\nabla F(.)$ is bounded, i.e. $\exists K > 0$ s.t. $\max_{x \in \mathbb{R}^n} \{\nabla F(x)\} < K$, then any trajectory along the dynamics will converge either to a fixed point or a limit circle.

Proof. Let $E(x) := \frac{1}{2} x_\alpha x^\alpha - v_\alpha x^\alpha - F(W^\alpha_{\ \beta} x^\beta + b^\alpha)$, then $\tau \mathrm{d}x / \mathrm{d}t = -\nabla E(x)$. The $-x$ term will dominate the $W^T \cdot \nabla F(W \cdot x + b)$ term for $\|x\| > K \|W\|$, thus converges. So E is a Lyapunov function of the dynamics. $\qquad \square$

Example 13. Let $F(x) := \sum_\alpha \int^{x^\alpha} \sigma(s) \mathrm{d}s$, where σ is sigmoid function. Then

$$\tau \frac{\mathrm{d}x}{\mathrm{d}t} = -x + W^T \cdot \sigma(W \cdot x + b) + v.$$

This coincides with the form in ref 1.

Example 14. Let $F(x) := \beta^{-1} \ln(\beta \sum_\alpha \mathrm{e}^{x^\alpha})$, $b = 0$, and $v = 0$, then

$$\tau \frac{\mathrm{d}x}{\mathrm{d}t} = -x + W^T \cdot \mathrm{softmax}(\beta W \cdot x).$$

This coincides with the form in ref 2.

Example 15. Let $v_i := W_{i,\cdot}$, i.e. the ith row of the matrix W. Assume $\|v_i\| = 1$ for $\forall i = 1, ..., n$. Let $F(x) := \beta^{-1} \ln(\beta \sum_\alpha \mathrm{e}^{x^\alpha})$, and $v = 0$, then

$$\tau \frac{\mathrm{d}x^\alpha}{\mathrm{d}t} = -x^\alpha + \sum_i p_i v_i^\alpha,$$

where $z_i := v_i \cdot x + b_i$ and then $p^i := \exp(\beta z^i) / \sum_j \exp(\beta z^j)$. The $\{(p_i, v_i) | i = 1, ..., n\}$ forms a categorical distribution.

Lemma 16. *Assume example 15. The Jacobian of the dynamics is*

$$J^{\alpha\beta}(x) = -\delta^{\alpha\beta} + \mathrm{Cov}_{p(x)}(v^\alpha, v^\beta),$$

where $\mathrm{Cov}_p(\cdot, \cdot)$ denotes the covariance given distribution p.

Proof. Directly,

$$
\begin{aligned}
J^{\alpha\beta} &\equiv \frac{\partial}{\partial x_\beta}\left(-x^\alpha + \sum_i v_i^\alpha p_i\right) \\
&= -\delta^{\alpha\beta} + \sum_{i,j} v_i^\alpha \frac{\partial p_i}{\partial z^j}\frac{\partial z^j}{\partial x_\beta} \\
&= -\delta^{\alpha\beta} + \sum_{i,j} v_i^\alpha v_j^\beta (p_i \delta_{i,j} - p_i p_j) \\
&= -\delta^{\alpha\beta} + \sum_i p_i v_i^\alpha v_i^\beta - \left(\sum_i p_i v_i^\alpha\right)\left(\sum_j p_j v_j^\beta\right) \\
&= -\delta^{\alpha\beta} + \mathbb{E}(v^\alpha v^\beta) - \mathbb{E}(v^\alpha)\mathbb{E}(v^\beta) \\
&= -\delta^{\alpha\beta} + \mathrm{Cov}_p(v^\alpha, v^\beta).
\end{aligned}
$$

And notice that the only variable that depends on x is p. So we insert x and gain the result. $\square$

For instance, at fixed point $x = v_1$, $p = (1, 0, ..., 0)$. $\mathrm{Cov}_{p(v_1)}(v^\alpha, v^\beta) = v_1^\alpha v_1^\beta - v_1^\alpha v_1^\beta = 0$. So $J^{\alpha\beta} = -\delta^{\alpha\beta}$ is negative defined, indicating that the fixed point is stable.

4 References

1. On autoencoder scoring.

2. Hopfield networks is All You Need.